In Databricks Serverless, we cant see the Spark UI, we can only see the performance.

Q1) Read data from ADLS in parquet and write in delta format along with table

Ensure to give same path

You can also write saveAsTable instead of save to do in a single step

Q2) Incremental Data Loading → use Autoloader

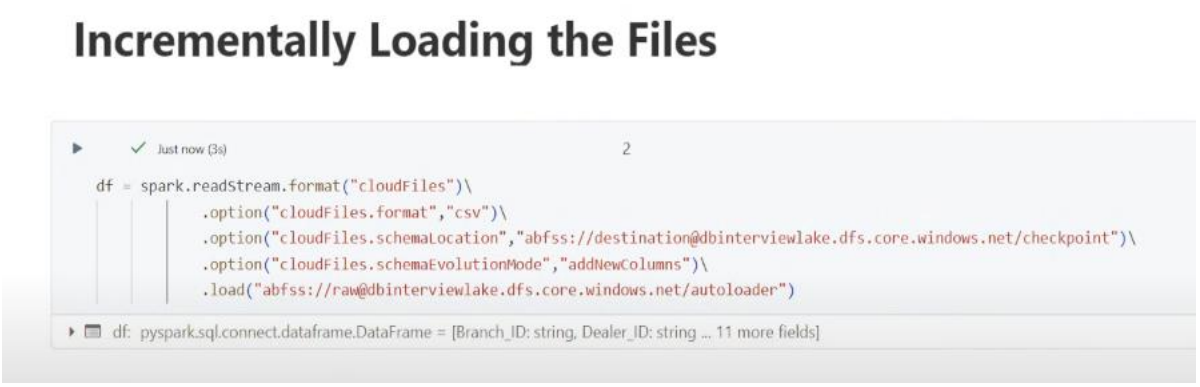

Schema Evolution happens in Autoloader by default, schemaEvolutionMode happens in Autoloader by default no need of explicitly mentioning it.

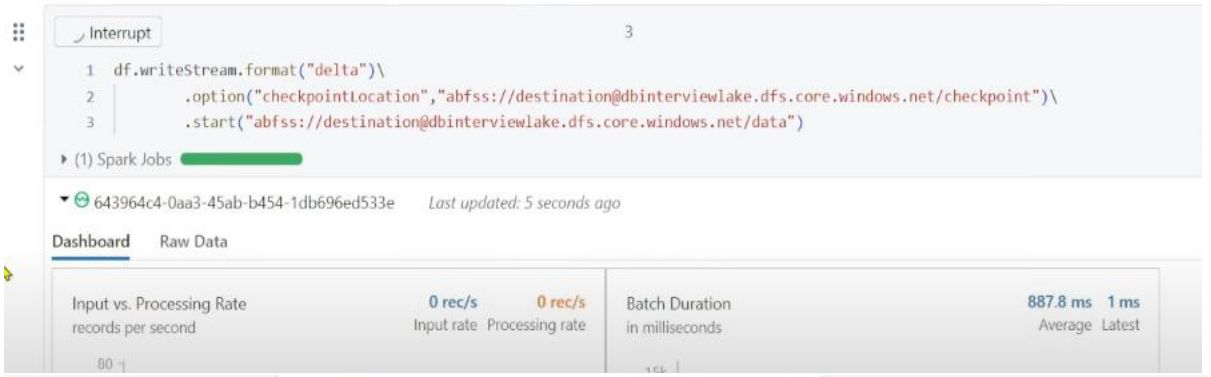

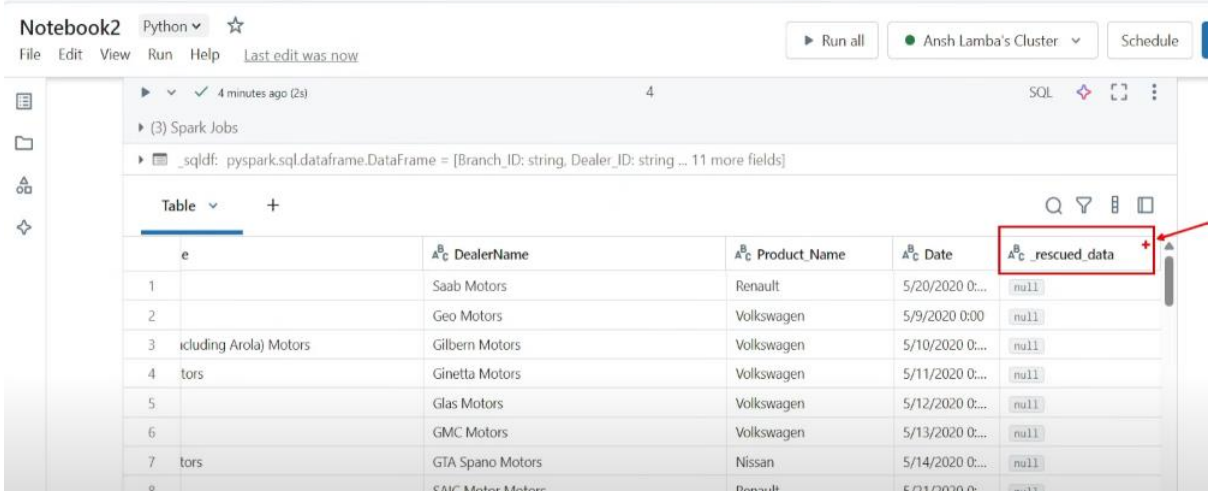

Schema is present in this location

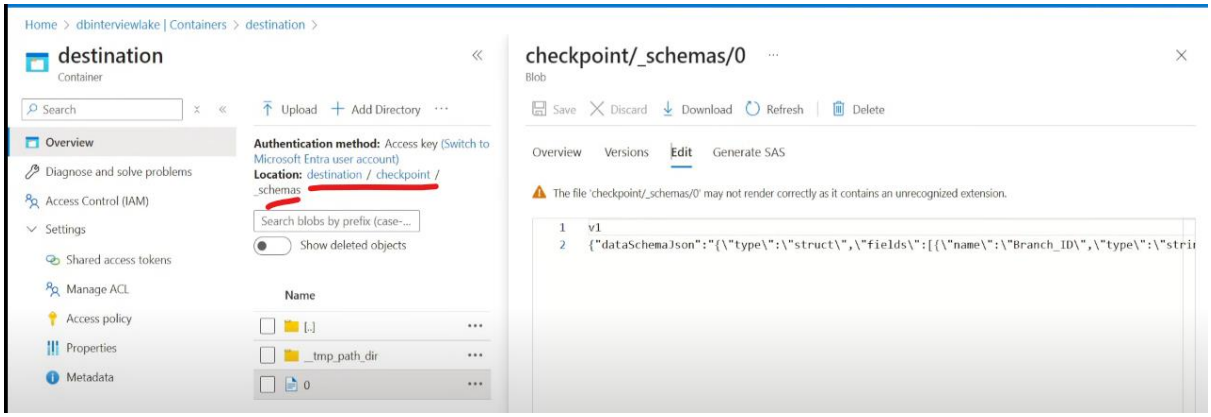

Now lets say there is schema change,and you want to rescue those columns, in destination I don't want to add any further columns but I want to store those columns

Incrementally Loading the Files

```
df = spark.readStream.format("cloudFiles")\
        .option("cloudFiles.format","csv")\
        .option("cloudFiles.schemaLocation","abfss://destination@dbinterviewlake.dfs.core.windows.net/checkpoint")\
        .option("cloudFiles.schemaEvolutionMode","rescue")\
        .load("abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader")
```

df: pyspark.sql.dataframe.DataFrame = [Branch_ID: string, Dealer_ID: string ... more fields]

All the new columns can be found in rescued_data column

How does Auto Loader schema evolution work?

Auto Loader detects the addition of new columns as it processes your data. When Auto Loader detects a new column, the stream stops with an `UnknownFieldException`. Before your stream throws this error, Auto Loader performs schema inference on the latest micro-batch of data and updates the schema location with the latest schema by merging new columns to the end of the schema. The data types of existing columns remain unchanged.

Databricks recommends configuring Auto Loader streams with Databricks Jobs to restart automatically after such schema changes.

Notebook2 Python

File Edit View Run Help Last edit was 6 minutes ago

Interrupt | Ansh Lamba's Cluster | Schedule

Table +

	Product_Name	Date	_rescued_data
1	Chevrolet	2/25/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
2	Mahindra	2/13/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
3	Chevrolet	2/26/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
4	Mahindra	2/14/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
5	Chevrolet	2/28/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
6	Mahindra	2/16/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
7	Chevrolet	2/29/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
8	Mahindra	2/17/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
9	Mahindra	2/18/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
10	Chevrolet	3/2/2020 0:00	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
11	Mahindra	2/19/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
12	Mahindra	2/20/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
13	Mahindra	2/21/2020 0:...	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir
14	GMC	3/5/2020 0:00	{"ReturnFlag":"1","_file_path":"abfss://raw@dbinterviewlake.dfs.core.windows.net/autoloader/raw_data_thir

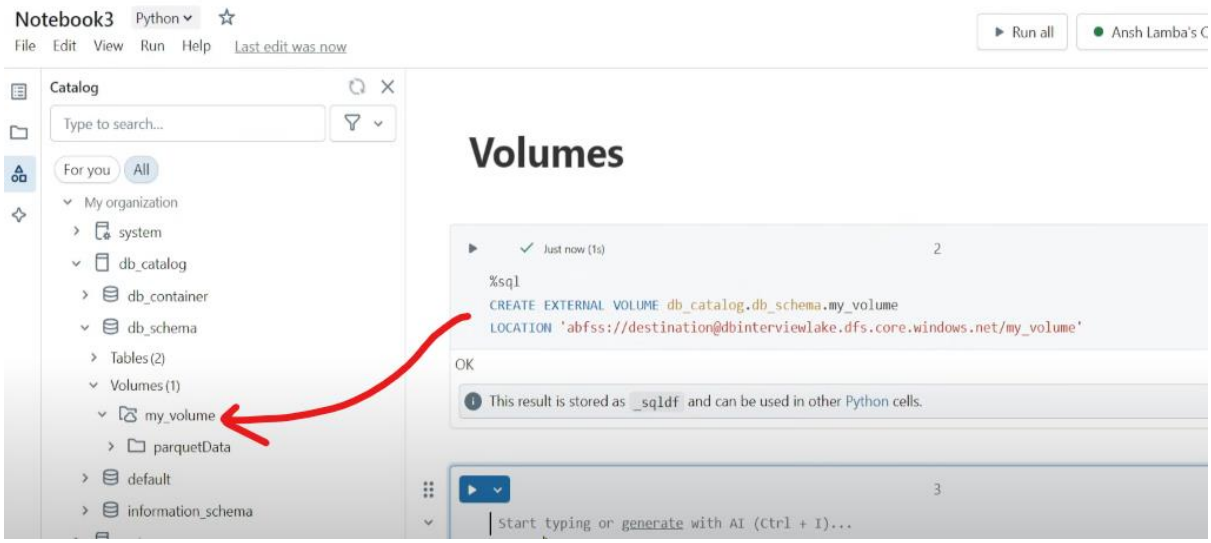

Volumes are like Files in Fabric LH

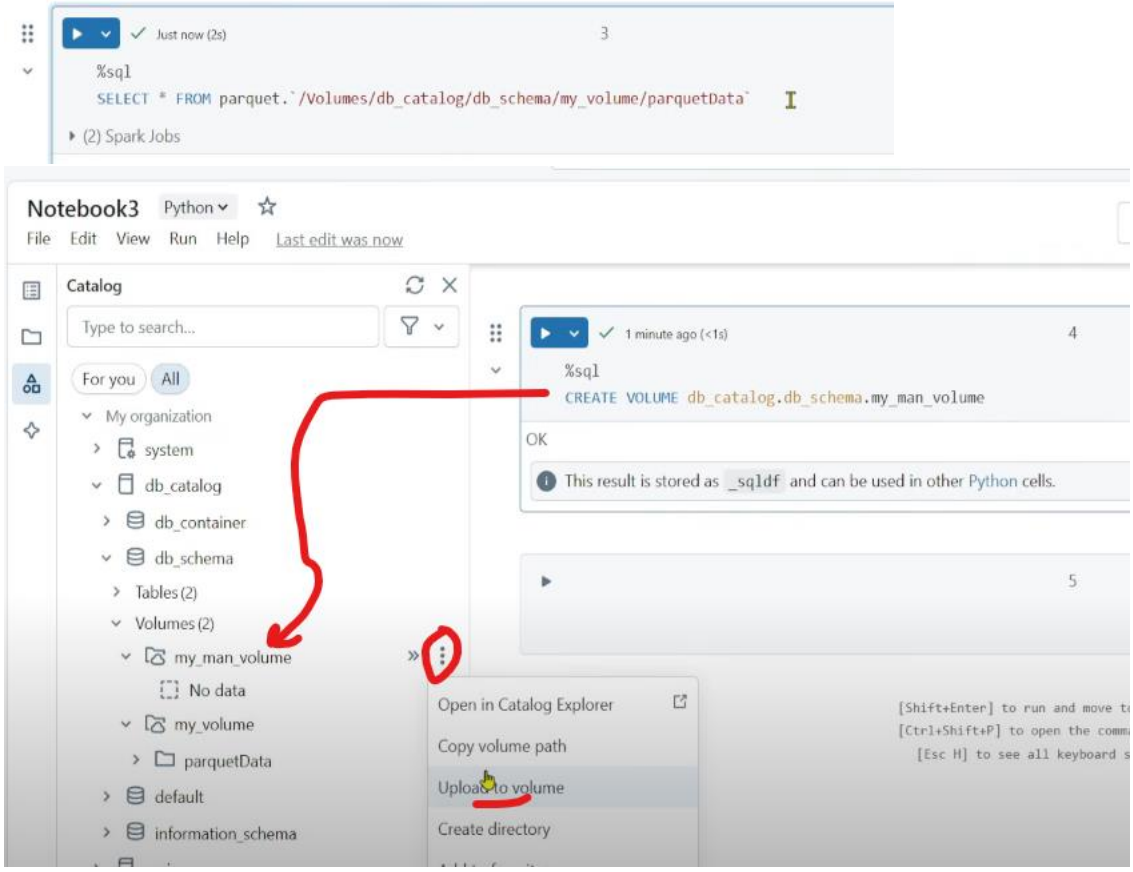

To insert data → Upload to Volume

Liquid Clustering

Liquid Clustering creates dynamic clusters on top of data, dynamic query behaviour

Use liquid clustering for Delta tables

Article • 03/28/2025 • 4 contributors

Feedback

In this article

Delta Lake liquid clustering replaces table partitioning and `ZORDER` to simplify data layout decisions and optimize query performance. Liquid clustering provides flexibility to redefine clustering keys without rewriting existing data, allowing data layout to evolve alongside analytic needs over time. Liquid clustering applies to both Streaming Tables and Materialized Views.

> **Important**
>
> Databricks recommends using Databricks Runtime 15.2 and above for all tables with liquid clustering enabled. Public preview support with limitations is available in Databricks Runtime 13.3 LTS and above.

What is liquid clustering used for?

Databricks recommends liquid clustering for all new Delta tables, which includes both Streaming Tables (STs) and Materialized Views (MVs). The following are examples of scenarios that benefit from clustering:

- Tables often filtered by high cardinality columns.
- Tables with significant skew in data distribution.
- Tables that grow quickly and require maintenance and tuning effort.
- Tables with concurrent write requirements.
- Tables with access patterns that change over time.
- Tables where a typical partition key could leave the table with too many or too few partitions.

SQL

SQL Copy

```sql
-- Create an empty table
CREATE TABLE table1(col0 int, col1 string) CLUSTER BY (col0);

-- Using a CTAS statement
CREATE EXTERNAL TABLE table2 CLUSTER BY (col0)  -- specify clustering after table name, no
LOCATION 'table_location'
AS SELECT * FROM table1;

-- Using a LIKE statement to copy configurations
CREATE TABLE table3 LIKE table1;
```

Python

```python
# Create an empty table
(DeltaTable.create()
  .tableName("table1")
  .addColumn("col0", dataType = "INT")
  .addColumn("col1", dataType = "STRING")
  .clusterBy("col0")
  .execute())

# Using a CTAS statement
df = spark.read.table("table1")
df.write.clusterBy("col0").saveAsTable("table2")

# CTAS using DataFrameWriterV2
df = spark.read.table("table1")
df.writeTo("table1").using("delta").clusterBy("col0").create()
```

You can enable liquid clustering on an existing unpartitioned Delta table using the following syntax:

```sql
ALTER TABLE <table_name>
CLUSTER BY (<clustering_columns>)
```

Enable or disable automatic clustering

To create a new table with automatic liquid clustering enabled, use the following syntax:

```sql
CREATE OR REPLACE TABLE table_name CLUSTER BY AUTO;
```

You can also enable automatic liquid clustering on an existing table, including tables that previously had manually specified keys, as shown in the following example:

```sql
ALTER TABLE table_name CLUSTER BY AUTO;
```

You can also alter tables with automatic liquid clustering enabled to use manually specified keys.